

ECS 189G

Homework 2

April 24, 2019

POS Tagger

This assignment is about POS tagging. You should be aware that POS tagging is widely used, and that it's widely seen in the NLP community as a solved problem (or at least well enough solved that most people don't put much effort into improving POS tagging for its own sake). It's a good problem for learning about NLP, though, and that's what we are aiming for in this exercise. For this assignment, you may only use the datasets provided. We realize that you could probably track down the test data we will evaluate your models on, but that would be cheating – so don't.

Download and unpack the tar file from canvas. There is a `README` file, three Perl scripts, one Python script, and some data. Using what we've given you, you can build a baseline bigram HMM POS tagger on the standard training set from the Penn Treebank (sections 2 -21, `ptb.2-21.*`) and evaluate it on the development set (section 22, `ptb.22.*`). (Actually, we already did that, but you should take two minutes to follow the `README` and do it yourself to make sure you understand the tools.) Note that we have not given you the standard test set output.

Task 1

A learning curve is a useful tool that lets you visualize how a model's performance depends on the amount of training data. You can vary the amount of training data by creating smaller sub-corpora from the original full corpus.

The learning curve plots a performance measure evaluated on a fixed test set (y-axis) against the training dataset size (x-axis). Generate a learning curve for the bigram HMM as we've provided it, using section 22 to evaluate. What are your thoughts about getting more POS-tagged data and how that would affect your system?

Task 2

This task is a prelude to the next task. Familiarize yourself with `viterbi.pl`. This is a bigram viterbi model written in perl. Your job is to implement this in python as `viterbi.py`. You do not need to follow the perl script, as long as your implementation is correct. This will be a good exercise in implementing the viterbi algorithm and to prepare a base model for the next task.

Task 3

Come up with a way to improve the model. Again, you can only train on the training data (no external resources are allowed). Some relatively easy ideas: implement a trigram HMM, or try to smooth the probability estimates. It may be helpful to know that many tagging schemes encode tags in such a way that if you replaced every tag by its first letter only, the tags would still be meaningful, only more coarse (this holds for all of the datasets used in this assignment).

You are constrained to using an HMM. You are most likely needing to replace `train_hmm.py` (or `train_hmm.pl`) according to the model of your choosing. Similarly, you are likely needing to change `viterbi.py` according to your model. Create new files for your new model. Describe your approach clearly. How well does your model perform on the development data? Run your tagger on `ptb.23.txt` (the test data) and turn in the output so we can evaluate it using `tag_acc.pl`.

Task 4

Train the baseline model and your model on the Japanese (`jv.*`) and Bulgarian (`btb.*`) training datasets, and compare them on the test sets. Report the performance in a readable way, and discuss. You should look at the data

(obviously!) and give some analysis – what factors lead to the differences among performance on English, Japanese, and Bulgarian? What about the baseline and your model makes them relatively better or worse on the data in these other two languages?

Task 5

Write a new evaluation script that calculates quantities that help elucidate the differences between the two models further, by looking at errors in a more fine-grained way. In your answer, explain what your script does and show its output.

Submission

Please submit two files to canvas. a .pdf containing your written work, and a .tgz archive containing the following items. Note that you are not just submitting code for this assignment; we rely on you to describe your approach well.

.pdf a pdf file containing your answers to all questions, scores, and the plot for task 1.

.tgz Your modified code (what we gave you + your modifications, compressed) the output of your new model on ptb.23.txt (we will evaluate it against gold-standard POS tags),

Code of Conduct

Please be reminded of the code of conduct as posted on the syllabus and on the OSSJA website. You must not share any parts of your code to others within and outside of the class. This includes, but not limited to, sharing screenshots of your code, posting or hosting your code onto platforms such as GitHub, obtain materials from previous versions of the class. At the same time, you must not use any code that is not written by you in this assignment. Your codes will be screened for plagiarism. We reserve the right to take disciplinary action to any offenders, which include reporting to Judicial Affairs.